

Principled Partisan Affect?

Evidence from Two Survey Experiments

Douglas J. Ahler

Carolyn E. Roush

Gaurav Sood

Abstract

Is partisan affect principled? People who value accountability should like partisans more when they withdraw support from a leader who has behaved badly, whether those partisans are allies or opponents. Two survey experiments examine whether portrayals of accountability and fair-mindedness change evaluations. In a 2013 MTurk study, respondents read that Democrats' support for Obama or Republicans' support for Cruz declined or remained high after a controversy. Average warmth toward fellow partisans was lower after declining support (difference: -4.3 points on a 0–100 scale; 95% CI [-8.9, 0.2]). This estimate points toward valuing loyalty over accountability and leaves little room for a positive reward for withdrawing support, although it does not establish a negative effect. For opponents, the estimated difference was small and positive (1.2; [-4.1, 6.5]). In a 2018 YouGov study, portraying opponents as fair-minded improved favorable trait ratings by an estimated 2.5 points ([-0.3, 5.4]), but warmth by only 0.5 ([-2.8, 3.9]). Portraying allies as biased produced no clear penalty; many readers also reported that their party was about as biased as expected. The evidence suggests a possible own-party loyalty response and modest benefits from portraying opponents as fair-minded, with little estimated change in warmth toward opponents. It does not establish a partisan double standard: the pilot's own-opposition difference is imprecise, and YouGov did not show the same behavior by both targets. Bundled article features, uncertainty, and missing warmth ratings also limit conclusions about why affect changes or how little it can change.

1 The question

Dislike of the other party need not arise solely from disagreement about policy. Partisans also judge the people on the other side, often unfavorably (Iyengar et al., 2012). They may dislike opponents partly because they see them as blindly partisan: willing to defend their leaders regardless of what those leaders do. If these judgments help sustain animosity, seeing opponents hold their leaders accountable should make them more likable. Applied impartially, the same principle would favor accountability among one's own allies and penalize their refusal to acknowledge wrongdoing. The question is whether partisan affect responds to evidence in this way.

Partisan loyalty could instead shape the response. A respondent might welcome opponents' willingness to abandon their leader while resenting fellow partisans who do the same. Observing that pattern would raise a further question: does the respondent apply different standards, or interpret the behavior differently depending on the party? Withdrawal of support is evidence

of principle only if the reader sees it as warranted. A controversy alone does not establish that respondents agree a leader behaved badly.

Evidence that people see more bias in others than in themselves (Pronin et al., 2002) does not resolve these questions. Nor does it establish that changing perceptions of bias changes affection. News portrayals can affect political judgments: Levendusky and Malhotra (2016), for example, find that coverage emphasizing partisan polarization increases dislike of the opposition. We examine whether portrayals of accountability and fair-mindedness move evaluations in the other direction, and whether portrayals of bias diminish affection for one’s own side.

The pilot’s estimate points toward preferring loyalty from one’s own side: reported withdrawal of support after controversy lowers estimated warmth toward fellow partisans. The later study suggests modest benefits from portraying opponents as fair-minded, but little estimated movement in warmth and no clear penalty for biased allies. These patterns warrant attention despite their uncertainty. Neither study, however, establishes that readers recognize the same conduct and then judge it differently across party lines.

2 Design and estimands

The studies differ in recruitment, stimuli, measurement, and the comparisons they permit. We analyze each separately. In both, the unit is the respondent, party leaners are included with their party, and the outcomes are measured immediately after the article. The reported intervals are two-sided 95% intervals. These are exploratory analyses of completed studies; we found no preregistered analysis plan in the study materials.

For each comparison, we estimate the treatment difference separately among Democrats and Republicans and average the differences using their shares in the study’s full eligible partisan sample. If p indexes party, s_p is that share, and a and b are the compared arms, the estimand is

$$\hat{\tau} = \sum_p s_p (\bar{Y}_{pa} - \bar{Y}_{pb}).$$

This prevents differences in the party composition of the realized arms from changing the pooled comparison. Each outcome uses the respondents who answered it; the standardization shares remain fixed across outcomes. Unweighted intervals use the sum of independent cell-mean variances and Welch–Satterthwaite degrees of freedom. The supplemental weighted estimates treat the supplied weights as fixed, use within-cell weighted means and HC3 variances, and standardize to weighted party shares. They address a different population composition from the unweighted estimates.

The questionnaires describe random allocation among eligible articles. The delivered assignments agree with their partisan routing. We do not have the randomizer’s execution logs. Causal readings require the stated randomization to have been implemented, the assigned stories to have been delivered, no interference between respondents, and an adequate account of outcome non-response. Random assignment identifies an effect of an article package. It does not identify a separate effect of perceived bias when other article features change too.

3 Declining support after controversy: MTurk, 2013

The first experiment was fielded on Amazon Mechanical Turk in November 2013. The source file contains 930 records, of which 919 have a known article assignment. The 11 records without assignment also lack the two party thermometers; we do not treat them as an assigned control group. Before treatment, an attention question instructed respondents to select only blue and yellow. Requiring those choices and no others excludes 39 records. The main sample consists of 724 assigned partisans passing this check: 552 Democrats and 172 Republicans. This convenience sample is not representative of American adults.

Respondents received one of five researcher-modified news stories: two about Democrats’ approval of Barack Obama after controversies surrounding his administration, two about Republicans’ approval of Ted Cruz after the government shutdown, and a control about support for Toronto mayor Rob Ford despite his drug scandal. Within each American party’s stories, approval either declined or remained high. “Sustained” describes the latter condition; its reported approval rose slightly rather than remaining numerically identical. Approval levels, trajectories, quotations, and accompanying narrative differed across the variants. We therefore compare versions of a story rather than treating the manipulation as a pure change in reasoning quality. Interpreting declining support as accountability additionally assumes that readers viewed the leader’s conduct as wrong and the withdrawal of support as justified. Those judgments were not independently randomized.

The main contrasts compare declining with sustained approval of the same leader. The outcome is a 0–100 feeling thermometer for Democrats or Republicans, oriented to the respondent’s own party or the opposition. Randomly ordered outcome blocks are combined only after checking that overlapping responses do not conflict. Every record with a known article assignment has both thermometers. Comparisons with the Ford control, shown in the appendix, additionally change the leader, party context, and controversy.

Declining support for one’s own leader did not produce a clear reward for holding that leader accountable. The estimated change in own-party warmth is -4.3 points $([-8.9, 0.2])$; Figure 1). The estimate points toward a preference for loyalty over accountability: respondents liked fellow partisans less when they read that support for their leader had declined. Its interval leaves little room for a positive reward for withdrawing support, but includes zero and therefore does not establish that declining support reduces warmth. For the opposition, the estimate is 1.2 $([-4.1, 6.5])$. This small positive estimate fits the hypothesized principled response: opponents may become more likable when they appear willing to hold their own leader accountable. The interval also permits a negative response or a larger positive one. The difference between these two effects is 5.5 points $([-1.4, 12.5])$. Thus the data do not establish that the response differs according to whether the story concerns allies or opponents.

The party-specific estimates are less precise. The opposition contrast is -0.8 among Democrats $([-6.5, 5.0])$ and 7.6 among Republicans $([-5.3, 20.5])$. The contrast between an imprecise positive estimate and an imprecise near-zero estimate is not itself evidence of a party difference. The own-party estimate remains negative under each attention rule; without the attention exclusion it is -3.1 $([-7.7, 1.4])$; Appendix Table 4). Its direction is stable across these choices, while its magnitude and precision vary. The pilot thus points toward a possible loyalty response to these stories, without




Figure 1: MTurk: declining rather than sustained approval after controversy. Points are differences in warmth toward the party described; bars are 95% Welch intervals. The pooled estimates standardize to the party composition of the 724 eligible respondents. Each contrast uses only the two relevant article arms.

identifying that motive or showing that respondents knowingly punish fair-mindedness.

4 Portraying allies and opponents: YouGov, 2018

YouGov fielded the second study from May 11 to May 29, 2018. Its codebook reports 1,596 interviews, from which 1,500 respondents were matched on gender, age, race, and education; subsequent weighting used propensity calibration and presidential vote. These are the vendor’s recruitment and adjustment procedures, not probability sampling. The delivered file contains 1,500 records. We report unweighted experimental comparisons first and weighted estimates as a sensitivity analysis. Interpreting either as an effect for all American adults requires assumptions about selection and treatment-effect generalization.

The party measure used before assignment identifies 641 Democrats and 492 Republicans, including leaners, for a main sample of 1,133 partisans. There are also 316 independents and 51 respondents selecting “not sure.” They have no defined own-versus-opposition party in our analysis and are reported separately in the appendix. We use the pre-assignment party measure to define the main sample and orient outcomes.

Table 1: The three arms available to each partisan group in YouGov.

Respondents	Control	Fair-minded opposition	Biased own party
Democrats	217	210	214
Republicans	166	164	162

The portrayed party depends on the respondent’s pre-assignment party. Democrats could read about fair-minded Republicans or biased Democrats; Republicans could read about fair-minded Democrats or biased Republicans.

Each partisan was eligible for a control article, an article portraying the opposition as fair-minded, or an article portraying their own party as biased (Table 1). The fair-minded articles described economic evaluations that remained stable across the transition from Obama to Trump and included favorable expert commentary. The biased articles described economic evaluations shifting after the election and included disparaging commentary from a political opponent. The control described aggregate economic views. The articles also differed in dates, framing, sources, and party references. These are experimental materials, not verified news reports or evidence that the depicted partisans actually reasoned well or poorly. Economic expectations can rationally change when the presidency changes; a change in expectations alone is not a sufficient test of bias.

Crucially, partisans were never assigned to fair-minded-own-party or biased-opposition articles. The design therefore cannot estimate the effect of the same portrayal for both partisan targets. Comparing the two available treatment effects changes both the target and the portrayal. Calling that difference a partisan double standard would exceed what was randomized.

4.1 Two distinct outcomes

The feeling thermometers rate the Democratic Party and Republican Party from 0 to 100. The pilot instead named Democrats and Republicans. We do not treat these as identical objects of evaluation or combine the studies' estimates.

The trait battery asks how well each of six descriptions fits each party: ignorant, sincere, open to reason, smug, patriotic, and hypocritical. Responses run from very poorly to very well. We reverse the three unfavorable descriptions and average available answers, rescaling the mean to 0–100. Higher values indicate more favorable judgments. This equally weighted index summarizes the measured descriptions; it is not a validated latent measure of fair-mindedness. A one-category shift on every constituent item would change the index by 25 points. The questionnaire includes a selfishness item that is absent from the delivered data, so it is not part of the index.

There are 132 missing opposition thermometers among the main sample, compared with 8 wholly missing opposition trait indices. Only 1,088 respondents answer all six opposition trait items. We report response counts by arm, complete-item estimates, and bounded-outcome sensitivity below. Conditioning on an observed outcome can compromise a randomized comparison if response depends on treatment and potential evaluations. Similar response rates alone would not resolve that concern.

4.2 Results

The estimated response to fair-minded opponents is favorable but modest, particularly for warmth. The trait estimate is 2.5 points (95% CI [-0.3, 5.4]), relative to a standardized control mean of 28.6. This is equivalent to 0.10 of a response category averaged across the constituent items. The warmth estimate is 0.5 points ([-2.8, 3.9]), relative to a control mean of 18.5. The warmth estimate is close to zero, whereas the trait estimate suggests a modest improvement in how opponents are described. Both intervals include zero. The difference in their point estimates does not itself establish that traits respond more than warmth, and the warmth interval permits an improvement of several points among observed respondents.




Figure 2: YouGov: responses to portrayals of the opposition and one’s own party. Differences from control, standardized to the eligible partisan composition. Bars are 95% Welch intervals. The two panels measure different judgments on a common 0–100 reporting scale.

Table 2: YouGov primary comparisons.

Article / rated target	Measure	Control	Difference [95% CI]	<i>n</i>
Fair-minded opposition	Warmth	18.5	0.5 [-2.8, 3.9]	667
Fair-minded opposition	Traits	28.6	2.5 [-0.3, 5.4]	752
Biased own party	Warmth	66.5	1.5 [-1.8, 4.9]	739
Biased own party	Traits	63.5	1.9 [-0.9, 4.6]	756

Control means and differences are standardized over pre-assignment party. *n* counts observed outcomes in the two compared arms, not the full three-arm sample. “Traits” is the available-item favorable index. All outcomes run from 0 to 100. Intervals are unadjusted for multiple comparisons.

The article portraying one’s own party as biased does not yield a clear penalty. Its estimates are 1.5 points for warmth ([-1.8, 4.9]) and 1.9 for traits ([-0.9, 4.6]). Their positive signs should not be turned into a demonstrated reward for bias: both intervals include negative and positive effects, and the article bundles several changes. None of the four main YouGov comparisons, or the two pilot comparisons, rejects a zero effect at the conventional 5% level after Holm adjustment within study. We report these adjustments to make the set of comparisons explicit, rather than treating a threshold crossing as the definition of what the studies teach.

4.3 How respondents interpreted the articles

The absence of an own-party penalty has more than one possible explanation. Readers might excuse their party’s bias, but they might also learn little they regard as new or credible. After the affect outcomes, YouGov asked readers of the partisan articles whether the portrayed party was more, less, or about as biased as they had previously expected (Table 3). Among readers of the biased-own-party article, 16% said their party was more biased than expected, while 46% said it

was about as biased. Among readers of the fair-minded-opposition article, 37% said the opposition was less biased than expected and 34% said it was about as biased.

Table 3: YouGov: reported bias relative to prior expectations.

Relative to expectations	Fair-minded opposition	Biased own party
More biased	10.4% (39)	16.2% (61)
Less biased	36.9% (138)	16.8% (63)
About as biased	34.5% (129)	46.3% (174)
Don’t know	17.6% (66)	19.9% (75)
Missing	0.5% (2)	0.8% (3)

Entries are unweighted percentages (counts) of all partisans assigned to each article arm, including unknown and missing responses. The question follows the affect outcomes. Control respondents received a different question; the two displayed arms portray different partisan targets.

These answers qualify what “exposure to bias” means here. Many readers did not report an adverse surprise about their own party, and only a subset reported a favorable surprise about opponents. The experiment therefore cannot be read as showing how much affection changes when everyone accepts new evidence of bias or fair-mindedness. The answers are retrospective reports, not measured changes from a pre-treatment belief, and the control arm has no equivalent question. They cannot distinguish prior expectations, disbelief, inattention, or reinterpretation, or establish that any of these explains the affect estimates. We do not condition the treatment comparisons on these post-treatment responses.

4.4 Sensitivity and missing outcomes

The opposition trait estimate remains small and positive under the weighting and completeness alternatives (Figure 3). Survey weighting raises the estimated opposition trait effect to 3.2 ([0.1, 6.3]). Requiring all six items gives 1.9 ([-1.0, 4.8]) without weights and 2.6 ([-0.5, 5.8]) with weights. The weighted warmth estimate is 2.1 ([-1.9, 6.1]). These choices change magnitude and precision without reversing the direction of the trait estimate. The evidence for a small improvement is suggestive; it should neither be declared conclusive from the weighted interval nor dismissed because other intervals include zero. The item estimates in the appendix also show why the composite should not be equated with a single judgment of hypocrisy.

The warmth result is particularly exposed to missing outcomes. Assigning every missing thermometer either endpoint of its scale gives worst-case bounds of [-10.5, 13.3] for the full-sample opposition contrast. Applying the same exercise separately to every missing constituent trait answer gives [1.3, 3.8] for the fixed six-item mean. These are ranges over possible completions of this dataset, not sampling confidence intervals. The positive lower trait bound describes the realized completed-data contrast; it does not remove randomization uncertainty or prove a positive population effect. The thermometer bounds are much wider and show that observed-case estimates alone do not sharply locate its full-sample effect.

No single observed respondent drives the available-item trait estimate: holding the party standardization fixed, deleting one respondent at a time yields estimates from [2.34, 2.69]. A permutation check that shuffles assignment within party among respondents with observed outcomes gives




Figure 3: YouGov opposition effects under weighting and item-completeness choices. Bars are 95% intervals: stratified Welch for unweighted estimates and fixed-weight HC3 for weighted estimates. Complete-item analysis also selects a different set of respondents. Weighted and unweighted estimates standardize to their respective eligible party shares.

$p = 0.084$ for this contrast. This check is conditional on outcome observation; it is not a remedy for treatment-dependent missingness.

5 What we learn about principled affect

An impartial accountability account predicts a reward for partisans who respond appropriately to misconduct, including those on one’s own side. The pilot’s own-party estimate points in the opposite direction. Its interval leaves little room for a positive reward, and the estimate remains negative under alternative attention rules. Yet a negative effect is not firmly established, and a preference for loyalty is one interpretation, not an identified motive. Respondents could have disagreed that withdrawal was justified, or reacted to other differences between the articles.

The opposition results fit a more encouraging but modest possibility: learning that opponents hold their leader accountable or reason fairly may weaken the impression that they are blindly partisan. Estimated warmth changes are 1.2 points in MTurk and 0.5 in YouGov; the later study’s trait estimate is 2.5 points. These point estimates suggest modest favorable responses, with little immediate softening of animosity. Their intervals include zero and permit larger improvements, however, and missing YouGov thermometers weaken any claim that the full-sample effect must be small. Different stories and objects of evaluation make this a resemblance across studies, not a pooled estimate or direct replication.

The absence of a clear penalty for biased allies does not establish indifference to bias. A story that confirms expectations need not change affection, even if those expectations already affect

how much a reader likes the party. The retrospective answers make this explanation plausible: many readers reported no surprise, and relatively few said their own party was more biased than expected. They do not establish the explanation. Likewise, the small opposition warmth effect cannot distinguish limited belief change from limited affective response to a belief change.

An apparent own-party loyalty response alongside favorable opposition estimates invites a double-standard interpretation. The pilot’s imprecise difference between these effects and YouGov’s missing treatment cells prevent that conclusion. A sharper test would independently randomize partisan target and evidence of accountability or bias, hold the surrounding story constant, and establish whether readers regard the conduct as warranted. Measuring prior expectations before treatment would help distinguish information that changes perceptions from information that confirms them. The present evidence points to possible selective responsiveness, with small estimated improvements in warmth toward opponents. Establishing its size, source, and persistence requires more informative treatments and more precise estimates.

Data and replication

The replication materials are available at <https://github.com/soodoku/principled>. They include deidentified respondent data, study materials, the complete R analysis, generated estimates, and this manuscript. Original source files are retained separately. `make check` reproduces the paper and runs linting and tests. The data documentation specifies exclusions, coding, missing values, inference, and the optional export from original files.

References

- Iyengar, S., Sood, G., and Lelkes, Y. (2012). Affect, not ideology: A social identity perspective on polarization. *Public Opinion Quarterly*, 76(3):405–431.
- Levendusky, M. and Malhotra, N. (2016). Does media coverage of partisan polarization affect political attitudes? *Political Communication*, 33(2):283–301.
- Pronin, E., Lin, D. Y., and Ross, L. (2002). The bias blind spot: Perceptions of bias in self versus others. *Personality and Social Psychology Bulletin*, 28(3):369–381.

A Supplementary comparisons

The tables below preserve the distinction between an eligible sample and the respondents contributing to a particular observed-outcome comparison. Party-specific estimates use their own group’s observed cell means; the pooled estimates in the text instead standardize to both parties’ shares.

Table 4: MTurk attention rules: declining minus sustained approval.

Sample	Own party	Opposition
Strict attention	-4.3 [-8.9, 0.2]	1.2 [-4.1, 6.5]
No attention exclusion	-3.1 [-7.7, 1.4]	1.4 [-3.9, 6.6]
Loose attention	-4.1 [-8.7, 0.5]	1.2 [-4.1, 6.5]

Entries are thermometer differences with 95% Welch intervals. Strict attention requires only the two requested colors; loose attention allows additional choices. All rules use known assignments and partisan identification. The party shares are recalculated for each eligible sample.

Table 5: MTurk stories compared with the Ford control, by respondent party.

Target	Decline versus Ford	Sustained versus Ford
<i>Democrat respondents</i>		
Own party	-5.3 [-10.6, -0.0]	-1.1 [-6.4, 4.2]
Opposition	0.4 [-4.7, 5.6]	1.2 [-4.7, 7.1]
<i>Republican respondents</i>		
Own party	-3.3 [-12.7, 6.1]	1.4 [-8.7, 11.5]
Opposition	6.6 [-5.7, 19.0]	-1.0 [-11.1, 9.2]

Entries are differences in thermometer ratings, with 95% Welch intervals. These comparisons change the leader, national and partisan context, and controversy, as well as the described approval pattern. They do not isolate approval decline.

Table 6: YouGov response counts: observed / assigned.

Target / measure	Control n	Fair-minded n	Biased n
Opposition warmth	343/383	324/374	334/376
Opposition traits	381/383	371/374	373/376
Opposition traits: six items	374/383	359/374	355/376
Own-party warmth	374/383	367/374	365/376
Own-party traits	381/383	371/374	375/376

Assigned counts include all eligible partisans. Trait indices use any available item unless marked “six items.” Each main effect compares one treatment arm with control. Respondents in the other arm do not contribute to that contrast’s reported n .

Table 7: YouGov effects by respondent party.

Article	Outcome	Democrats	Republicans
Fair-minded opposition	Opposition warmth	0.9 [-3.6, 5.5]	-0.0 [-5.0, 4.9]
Fair-minded opposition	Opposition traits	2.1 [-1.7, 5.8]	3.1 [-1.3, 7.5]
Fair-minded opposition	Own-party warmth	-2.3 [-6.7, 2.2]	6.0 [0.7, 11.2]
Fair-minded opposition	Own-party traits	0.3 [-3.3, 3.8]	3.3 [-0.8, 7.3]
Fair-minded opposition	Warmth gap	-2.6 [-9.5, 4.2]	4.0 [-3.4, 11.4]
Biased own party	Opposition warmth	-1.9 [-6.3, 2.5]	3.2 [-2.0, 8.4]
Biased own party	Opposition traits	-0.2 [-3.8, 3.4]	0.9 [-3.3, 5.2]
Biased own party	Own-party warmth	-0.5 [-4.8, 3.9]	4.2 [-1.1, 9.4]
Biased own party	Own-party traits	-0.2 [-3.8, 3.5]	4.5 [0.3, 8.6]
Biased own party	Warmth gap	2.6 [-4.4, 9.6]	-0.1 [-7.3, 7.1]

Unweighted differences from control with 95% Welch intervals. Outcomes run from 0 to 100 except the warmth gap, defined as own-party minus opposition warmth, which runs from -100 to 100 . Effects on the party not described and on the gap are secondary comparisons. Comparing the significance of two estimates is not a test of their difference.




Figure 4: YouGov constituent traits. Unweighted differences from control; 95% Welch intervals. Unfavorable descriptions are reversed so that positive values always mean more favorable ratings. Intervals are not multiplicity-adjusted; the replication output also reports Holm-adjusted p -values across the item comparisons.

Table 8: YouGov respondents without a partisan side: evaluations of Republicans.

Respondents	Article	Republican traits	Republican warmth
Independent	unbiased	5.6 [0.3, 10.9]	4.8 [-2.3, 11.9]
Independent	biased	2.3 [-3.6, 8.2]	2.0 [-5.5, 9.4]
Not sure	unbiased	-13.2 [-25.3, -1.1]	-3.0 [-24.3, 18.3]
Not sure	biased	-4.0 [-12.5, 4.6]	-7.6 [-27.7, 12.5]

Unweighted differences from control with 95% Welch intervals. “Unbiased” is the fair-minded-Republican article, and “biased” the biased-Republican article. Independents and respondents answering “not sure” to pre-assignment party identification are separate groups. The observed routing for “not sure” matches that for independents, but the questionnaire does not explicitly document that routing. These estimates concern Republican evaluations, not out-party affect.

B Computation and scope of uncertainty

For the primary unweighted estimates, let s_{pa}^2 and n_{pa} be the sample variance and observed count in party–arm cell (p, a) , and let c_a encode the arm contrast. The estimated variance is

$$\hat{V}(\hat{\tau}) = \sum_{p,a} s_p^2 c_a^2 \frac{s_{pa}^2}{n_{pa}}.$$

Writing each summand as v_{pa} , the Welch degrees of freedom are $(\sum v_{pa})^2 / \sum [v_{pa}^2 / (n_{pa} - 1)]$. This calculation conditions on the standardization shares and treats respondents as independent. It does not include uncertainty from sample recruitment or the selection of these particular stimuli.

For a weighted cell, normalized weights are $h_i = w_i / \sum_i w_i$. Its HC3 variance is $\sum_i \{h_i(Y_i - \bar{Y}_w) / (1 - h_i)\}^2$. We combine these cell variances with the squared weighted party shares and contrast coefficients, using observed comparison size minus the number of compared cells as the reference t degrees of freedom. The supplied weights’ estimation uncertainty is not included. No geographic or household cluster identifiers needed for a different dependence model are available in the public analysis data.

The replication tables also retain raw cell means, effective cell sample sizes, response-rate contrasts, all main outcome comparisons under both weighting choices, leave-one-item-out indices, exact case-deletion ranges, and the seeded permutation checks. These checks probe the reported estimator and selected analytic choices. They do not compensate for absent treatment cells, validate the news stories, or establish an unmeasured mechanism.