

Is Cynicism Taken as Evidence of Political Sophistication?

Gaurav Sood

Alexander G. Theodoridis

Abstract

The more cynical of two statements about politics receives lower sophistication ratings. A tweet saying politics attracts bad people averages 47.5 on a 0–100 scale, compared with 53.6 for a tweet saying politics attracts good people. The difference is -6.2 points (95% confidence interval [-9.4, -3.0]) among 711 respondents. Both tweets say the system rewards bad behavior. The result therefore concerns cynicism about who enters politics, holding criticism of the system constant. The comparison remains negative with covariate adjustment and survey weighting. A causal interpretation requires random tweet assignment within the rating sample; assignment fields are present, but the original questionnaire and randomization program are unavailable.

1 Question and comparison

Does expressing a darker view of politics make someone seem more sophisticated? Cynical targets have been judged better at cognitive tasks than less cynical targets, even though separate assessments of actual competence did not support that perceived advantage ([Stavrova and Ehlebracht, 2019](#)). This distinction motivates a question about reputation: does a cynical political statement make its speaker seem sophisticated? General competence judgments and political sophistication ratings need not respond to the same cues.

This study compares two statements that agree about the incentives in politics but disagree about the people it attracts:

Politics attracts good people but the system rewards bad behavior.

Politics attracts bad people and the system rewards bad behavior.

The images present the statements as tweets with the same visible account format and timestamp. The treatment changes “good people” to “bad people” and “but” to “and.” It thus compares these two complete statements, not a single word in isolation. The first statement may communicate nuance or make a distinction between people and institutions that respondents value. Agreement, nuance, and moral pessimism are plausible explanations for the rating difference; the design does not distinguish them.

The measured outcome is perceived sophistication. It is not a knowledge test, and the files do not establish the rated speaker’s actual knowledge. Nor is the comparison between cynicism and an optimistic view of politics: both statements criticize the system. A positive difference would support a sophistication premium for the more cynical statement in this setting. The observed difference is negative.

2 Data and research design

The supplied files cover the pre-election portion of the 2020 Cooperative Election Study UCM module. The larger file contains 1427 respondents, fielded September 30 through November 2. The other contains 1000 respondents, fielded October 3 through November 2. All respondents in the latter appear in the former, with identical treatment, order, outcome, and team-weight fields. We call this the matched subset, following the companion codebook’s description of the larger file as unmatched. General CES procedures use matching and post-stratification, but the module-specific selection and calibration instructions are not supplied.¹

The order field describes a split between rating the tweet before the election and reporting one’s own good/bad-people belief afterward, or doing those tasks in the opposite waves. The available outcomes implement only the pre-election part of that plan. There are 711 sophistication ratings and 716 belief responses, with no respondent observed on both. Other post-election variables exist, but neither of these questions has a post-election counterpart in the supplied files. The two files must not be stacked as independent samples.

Pre-election question	Tweet assignment	Assigned	Ratings	Beliefs
Sophistication	Good people	370	370	0
Sophistication	Bad people	341	341	0
Own belief	Good people	345	0	345
Own belief	Bad people	371	0	371

Table 1: Question routing in the full file. Tweet assignment is recorded for both routes. Only the sophistication route provides a rating outcome; the belief route is not another rating sample. Every respondent answers the question assigned to their pre-election route.

The treatment labels identify bad/bad, good/bad, and good statements. Only the first two have observations. Their names correspond to the supplied `ps_bb_tweet.jpg` and `ps_gb_tweet.jpg` images. The repository also contains other candidate stimuli, whose use cannot be established. Filename and label correspondence supports the two-image mapping; a survey program or display log would be needed to verify it independently. The codebook contains a short outcome description and response labels, not the full question stem. No preregistration is included, so the analysis choices below are retrospective.

2.1 Estimation and assumptions

The outcome runs from extremely sophisticated to extremely unsophisticated in seven ordered categories. We code it as $100(7-r)/6$, where r is the original response code. Extremely sophisticated is 100, the neutral response is 50, and extremely unsophisticated is 0. Neutral answers are retained. “Not asked” codes remain missing. All routed rating respondents supply a valid rating, so there is no observed item nonresponse in this sample; dropout before inclusion in the files cannot be assessed.

The primary estimate is the bad-people mean minus the good-people mean in the full rating sample. Its standard error is $\sqrt{s_1^2/n_1 + s_0^2/n_0}$, with a Welch–Satterthwaite t interval. The comparison

¹The [official CES FAQ](#) describes recruitment, matching, and weighting. It provides context, not a record of this module’s exact administration.

estimates an average assignment effect if tweet assignment is random within the rating sample, treatment is delivered as assigned, and there is no interference. The recorded split fields and counts are consistent with a survey experiment, but counts and covariate balance cannot verify randomization. Without that assumption, the estimate is an arm difference.

We report adjustment for age, gender, college education, the White race category, and five-category party identification. Each covariate is centered at its full rating-sample mean and interacted with treatment. The treatment coefficient therefore evaluates the contrast at the sample’s covariate distribution. These demographic and common-content fields are treated as pre-treatment; their timing should be confirmed against the questionnaire. All are complete in the rating sample. HC3 standard errors and residual-df t intervals accompany adjusted estimates. Balance statistics describe assignment groups rather than determine the adjustment set; the largest absolute standardized difference is 0.15.

The matched subset is analyzed both unweighted and with the supplied `teamweight`. The weighted contrast subtracts separately normalized weighted arm means and uses HC3 uncertainty with weights treated as fixed. It does not incorporate uncertainty from matching or estimating the weights. The weighted target differs from the full-sample target, and population generalization additionally requires valid calibration and transport of the contrast. Weight-based effective sample sizes are 167.5 and 160.9 in the good- and bad-people arms.

Respondents are the independent units under the assumed assignment mechanism. There are no repeated ratings to cluster. The two tweets are fixed stimuli; respondent-level uncertainty does not measure variation across possible tweets, speakers, or political contexts.

3 Estimates

The good-people tweet averages 53.6 and the bad-people tweet 47.5, yielding a difference of -6.2 points (95% interval [-9.4, -3.0]; SE 1.64; two-sided $p = 0.00018$). This is -0.37 of one response-category step. The point estimate supports a penalty for the more cynical statement, and the interval excludes a positive mean premium under the stated inferential assumptions.

The comparison arm is substantively informative. For the good-people tweet, 37.8% choose one of the three sophisticated categories, 41.4% choose neutral, and 20.8% choose an unsophisticated category. The corresponding percentages for the bad-people tweet are 24.9%, 43.1%, and 32.0%. Thus the result is not that one statement receives uniformly high ratings. Neutral is the largest single response category in each arm; the distributions differ on both sides of it. The scale midpoint is a response label, not a no-tweet control.

The adjusted difference is -5.9 points ([-9.2, -2.6]). Among the 505 rating respondents in the matched subset, the unweighted difference is -7.1 ([-10.9, -3.3]); applying team weights gives -7.9 ([-12.5, -3.3]). Inclusion in the matched subset is 70.5% in the good-people arm and 71.6% in the bad-people arm. Similar retention rates do not establish selection on outcomes is absent. The 206 respondents excluded from the subset have an estimated contrast of -3.9 ([-10.0, 2.1]), which is less precise and compatible with both a negative difference and a modest premium. The subgroup interval does not establish that exclusion changes the effect.

Greater cynicism receives lower sophistication ratings

Both tweets say the system rewards bad behavior




Figure 1: Perceived sophistication by tweet in the full rating sample. Means and 95% t intervals use all valid responses in each arm. The dashed line marks the neutral response. The estimate of the between-arm contrast uses its own standard error and does not rely on overlap of these arm intervals.

Analysis	<i>n</i>	Difference	95% interval
Full sample	711	-6.2	[-9.4, -3.0]
Covariate adjusted	711	-5.9	[-9.2, -2.6]
Matched subset	505	-7.1	[-10.9, -3.3]
Matched weighted	505	-7.9	[-12.5, -3.3]
Rounded scale	711	-6.2	[-9.4, -2.9]
Unmatched remainder	206	-3.9	[-10.0, 2.1]

Table 2: Bad-people minus good-people ratings, in 0–100 scale points. Intervals use Welch inference except for covariate-adjusted and weighted estimates, which use HC3 t inference. The rounded scale uses 0, 16, 33, 50, 66, 83, and 100. Sample restrictions and weighting change the target population; the rows are not independent replications.

3.1 Scale and influence checks

Equal spacing is a convention for the rating scale. A binary comparison that only asks whether a response is at least somewhat sophisticated gives a -12.9 percentage-point difference (95% interval [-19.7, -6.1]). The six cumulative thresholds all have negative point estimates, although several tail thresholds are imprecise (Table 3). The at-least-somewhat comparison has Holm-adjusted $p = 0.0012$ within that six-test exploratory family. This pattern is not created solely by the equal-spacing convention, but it does not establish population stochastic dominance at every threshold.

Both sample medians are neutral. The probability that a randomly selected bad-people rating exceeds a randomly selected good-people rating, counting a tie as one half, is 41.5%. A supplementary rank-sum test with tie adjustment and continuity correction gives $p = 0.00004$. This describes the rating distributions, not the probability that an individual’s rating would fall under treatment.

A two-sided permutation check gives $p = 0.00015$ using 19999 draws, seed 20261004, and the finite-draw correction $(k + 1)/(B + 1)$. Its approximate Monte Carlo standard error is 0.00009. This test conditions on the observed arm sizes and assumes complete random assignment within the rating route. The missing assignment program prevents verification that this was the actual assignment mechanism.

Deleting each respondent in turn leaves the main contrast between -6.3 and -6.1 points. Tables of leverage, design-matrix rank, adjusted and weighted influence, and covariate balance are generated with the analysis. The bounded outcome and complete item response limit the scope for extreme-value and missing-outcome artifacts within the observed rating sample.

4 Interpretation

These estimates contradict a sophistication premium for the more cynical of the two statements. Greater negativity about the people attracted to politics is associated with lower perceived sophistication, and under random assignment causes lower ratings in this sample. The direction persists across the full sample, covariate adjustment, the supplied subset, and weighting. The evidence supports that narrow finding, not a claim that every expression of cynicism is penalized.

The result also differs from the perceived cognitive advantage of cynical targets in [Stavrova and Ehlebracht \(2019\)](#). It is not a direct replication: the present comparison uses political statements that share criticism of incentives, rather than contrasting generally cynical and noncynical people. It therefore identifies a possible boundary to a broad sophistication-premium claim, without overturning that study’s evidence about other judgments and stimuli.

The separate belief sample provides context: 55.9% select the good-people statement and 44.1% the bad-people statement (the latter’s 95% interval is [40.5, 47.8]). Because both choices criticize political incentives, neither percentage measures how many people are cynical in a general sense. More respondents endorse the statement that receives higher sophistication ratings in the other sample. Agreement with the message could help explain the rating gap, but no respondent has both outcomes, so these files cannot test that mechanism or whether a person’s own cynicism moderates the contrast.

Exploratory treatment-interaction tests for gender, college education, age, and Republican identity are imprecise (Table 4). Their intervals admit substantively meaningful differences, so they do not establish uniform effects. A direct test of the agreement mechanism would require pre-treatment

beliefs and subsequent ratings for the same respondents, with the belief measurement timed to avoid priming the rating.

The main remaining design uncertainty is documentary: the full question stem, randomization instructions, and verified stimulus display are absent. Those materials would establish what was implemented. Additional stimuli varying cynicism and perceived nuance separately would be needed to determine why these statements differ and whether the result generalizes beyond them.

References

Olga Stavrova and Daniel Ehlebracht. The cynical genius illusion: Exploring and debunking lay beliefs about cynicism and competence. *Personality and Social Psychology Bulletin*, 45(2):254–269, 2019. doi: 10.1177/0146167218783195.

A Response distributions and additional estimates

Ratings shift away from sophisticated




Figure 2: All seven rating categories, expressed as percentages within each tweet arm. The two panels use the same horizontal scale. Neutral responses remain in the denominator.

B Thresholds and interactions

Sophistication threshold	Difference	95% interval	Holm p
Extremely	-0.8	[-3.7, 2.1]	0.5804
Very or extremely	-3.5	[-8.3, 1.2]	0.4259
At least somewhat	-12.9	[-19.7, -6.1]	0.0012
At least neutral	-11.2	[-17.6, -4.7]	0.0038
Above very unsophisticated	-6.9	[-11.7, -2.1]	0.0202
Above extremely unsophisticated	-1.8	[-5.1, 1.5]	0.5505

Table 3: Cumulative rating contrasts in percentage points, bad people minus good people. All use the full rating sample and Welch intervals. Holm adjustment applies to the six p -values; the displayed intervals are pointwise, not simultaneous. “Above” excludes the named response and all less sophisticated responses.

Difference in tweet contrast	Estimate	95% interval	Holm p
Female minus male	-0.1	[-6.7, 6.6]	1.00
College minus no college	0.3	[-6.4, 7.1]	1.00
Age 50+ minus under 50	-1.1	[-7.5, 5.4]	1.00
Republican minus other	3.7	[-3.9, 11.3]	1.00

Table 4: Exploratory differences in the tweet contrast, in 0–100 scale points. Each row is a separate treatment-by-group interaction with HC3 t inference. College means a four-year or postgraduate degree. Republican means the Republican response in the five-category party item, without reallocating leaners. Holm adjustment applies to these four tests. Intervals are pointwise.

C Reproduction

Run `make restore` and `make check` from the repository root. The pipeline reads the supplied SPSS files, validates their overlap and question routing, writes aggregate tables, generates figures, compiles this report, and runs lint and tests. The dependency lockfile records the R packages. GitHub Actions runs the same checks. Numerical prose is generated through `tabs/macros.tex`; tables and figures use the same analysis outputs. No live data download is needed.

The tests independently reproduce Welch inference, derive weighted HC3 uncertainty from arm-level residuals, verify coding and denominators, test the permutation correction, compare influence diagnostics with actual deletions, and check recovery of known effects in simulated data. The data guide and design note document every analysis variable and the interpretation assumptions.